CASE STUDY

When Site Issues Keep Repeating Because Nobody Owns the Full Story

Dubai Facilities Management Portfolio

We built a voice and automation layer that turns the first sign of a problem on a call into a trackable action, keeps it red and noisy until it is actually addressed, and stops "we will look into it" from dying inside Teams threads, call summaries, or someone's memory.

Build posture: If a complaint cannot be tied to a specific site, asset, vendor, and open work order, we do not label it "handled." We keep it red, assign an owner, and keep notifying until the proof points show it is truly addressed.

Client Snapshot

A Dubai-based facilities management operator running end to end upkeep for multiple corporate hubs owned by one real estate developer.

They cover hard FM and soft FM at the same time, so a normal week includes:

- HVAC hot spots
- Lift downtime
- Fire panel alarms
- Washroom complaints
- Cleaning quality
- Access issues
- Vendor coordination
- Penalty exposure when SLAs slip

They already had the basics. Calls were recorded. Transcripts and summaries existed. Tickets existed. Work orders existed.

> The most dangerous information lived in conversations, not in the systems that drive action. The site team heard it, nodded, and moved on. The same issue came back next week — and now it was not "a problem." It was a **trust problem.**

Where the Problem Showed Up in Real Life

This is what we anchored the build around. Not theory, not dashboards.

Scene 1 — Weekly Site Review Call

Someone says, "This is the third time this month we've had hot spots on Level 14." The site lead replies, "We closed the work order last week." Everyone moves on. Two days later, a tenant complains again and the developer's team asks, "Why are you closing tickets without fixing the root cause?" The FM operator now scrambles across CMMS history, vendor notes, and call clips to explain what happened.

Scene 2 — Lift Complaints That Look Like Isolated Tickets

The ticketing system shows two separate lift incidents across two towers. On calls, you can hear the pattern: "it keeps getting stuck around the same time every morning," "the vendor response is slow," "your security team gets shouted at." Nobody links the call evidence to the vendor and asset history early enough to stop the repeat.

Scene 3 — Soft FM Issues That Turn Into Politics

Cleaning supervisor says "team was there," tenant rep says "washroom still dirty," and the account team spends the day forwarding photos. The next call has the same sentence again: "This is still happening." The problem is no longer cleaning. The problem is the client believing you do not close loops.

Scene 4 — Commercial Pressure in Casual Language

A developer-side stakeholder says, "We'll apply penalties if this repeats," or "we're reviewing scope next quarter." In the moment it sounds like frustration. In reality it is a contract signal. If it is not captured as a red flag with an owner and a remediation plan, it becomes a surprise in the monthly steering call.

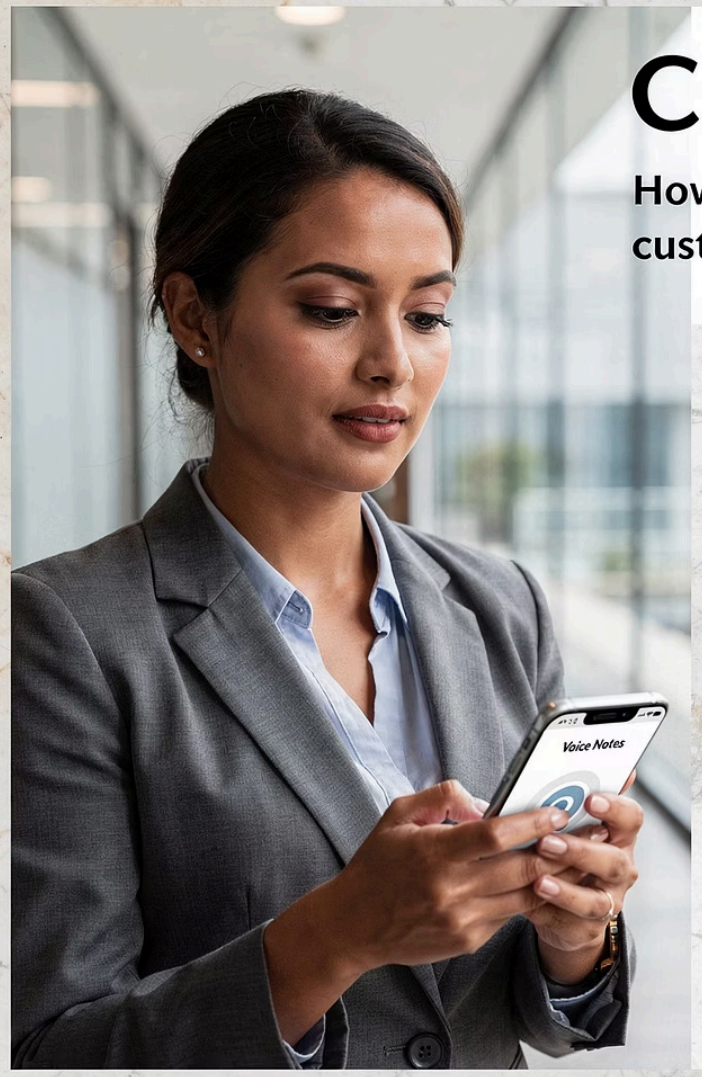

Case

How we tra
customer c

Project

Lorem we
abornverie
dust vrieet
throuk abc

Results

Lorem we
alaoptioms
ainglo and
nooigrathu

Proces

Lorem we
areavicen
comunicat
antlecatt h

What We Delivered

We shipped an **Operational Action Voice Copilot** with automation, designed to do one job consistently: the moment a risk shows up in a call, it becomes a structured action with owners and proof requirements, and it stays visible until those proof points exist.

Where It Sits in Their Day

We did not add a new portal and hope people change behaviour. We placed it inside the surfaces they already use.

1

Mobile Voice App

For site leads and portfolio heads — capture the real issue while it is fresh, not at end of day.

2

Account Review Dashboard

A "current red flags" panel and a weekly action packet used before the developer steering call.

3

Teams Notifications

Owners get notified the same day. If an item stays unresolved, reminders get louder and more specific.

What It Integrates With

- Call recording and transcript system already used by leadership
- CMMS for work orders, assets, vendor assignments, closure notes
- Ticketing for escalations and cross-team handoffs
- SLA and performance reporting
- ERP or invoicing workflows where penalties, credits, or scope changes are tracked
- Teams and email for routing and approvals

How It Behaves Day to Day

This is the part that mattered to operations leaders. They did not want another "insight tool." They wanted a system that **forces action.**

Turning a Call Moment Into a Red Flag That Cannot Be Ignored

A portfolio head speaks into the app right after a call: *"Hub A, Tower B, lift complaints again, vendor response time mentioned, developer asked why it keeps repeating."*

The copilot links the call moment, ties it to the site and asset list, creates the red flag, assigns the owner, and creates the remediation task trail inside CMMS and ticketing.

When Someone Asks for Status, the Answer Is Not a Story

A manager asks: *"What are the top repeat issues for Hub A this month and what is actually done, not promised?"*

The system returns a short list with proof points — open work orders, closure evidence, repeat complaint signals, and the current owner status. If the evidence is missing, it says that plainly and **keeps the item red.**

When It Is Not Sure, It Does Not Turn Vague Language Into a False Conclusion

If a call says "west wing AC is bad again" but the system cannot map it to a specific hub or floor, it asks one detail that a human can answer in two seconds. If it is still unclear, it creates a "needs clarification" task for the site lead, attaches the voice clip and transcript snippet, and notifies the owner so the question is resolved, not forgotten.

Red Stays Red Until the Proof Points Exist

A Red Flag Only Clears When the System Sees More Than One Proof Point — Not Just "Ticket Closed"

- The relevant work order is closed in CMMS with a meaningful closure note
- The same complaint does not reappear in calls in the next defined window
- SLA related items have the expected evidence — response timestamps or vendor attendance logs

If a manager wants to override and mark it addressed, they can — but the override is recorded and the system still watches for repeat signals.

The System Nags the Right People, Not Everyone

- If the issue is HVAC, the HVAC owner and the site lead get notified
- If it is cleaning quality, the soft FM lead gets notified
- If it touches penalties or credits, the commercial owner gets notified

Notifications are tied to what each user asked previously, so a portfolio head who asked *"why is Hub C messy every week"* will get proactive alerts when the same pattern starts forming again.

What Became Operationally Different

1

Problems Became Owned

Before this, problems were "known" but not owned. After this, the moment a complaint was heard, it became something you could point to, assign, and track without relying on someone's memory.

2

Site Reviews Had Teeth

Site reviews stopped being a conversation that disappears after the call, because each repeat issue was captured as a red flag with a clear owner and a due date.

3

CMMS Connected to Reality

CMMS and ticketing stopped being disconnected from what the client actually hears on calls, because the call evidence was linked to the work order and the vendor.

4

Leadership Walked In Prepared

Leadership stopped walking into the steering call with general updates, because the dashboard always showed what is still red and why it is still red.

What Surprised Us

> The biggest time sink was not fixing issues. It was **proving what was already done**, because the evidence lived in five places and still did not tell one clean story.

> The word **"again"** on calls was the earliest signal of future escalation, and it was being treated like emotion instead of a repeat-failure pattern.

> **Closing a ticket without closing the loop** creates more damage than leaving it open, because it trains the client to stop believing your system.

Does This Sound Familiar?

Do your weekly site calls include the same complaints, but your action system treats them like fresh incidents every time?

Do site leads say "it's closed" while the developer side says "it's still happening"?

Do issues get discussed in Teams calls and then disappear into summaries that nobody turns into work orders, owner tasks, and follow-through?

Do penalties and scope review threats show up in casual call sentences, and then surprise you later because nobody captured them as red flags?

If this sounds familiar, talk to us. We will map one hub, plug into your call recordings, CMMS, and ticketing, and ship a voice and automation layer that turns the first mention into action, keeps it visible until proof exists, and stops repeat issues from hiding inside conversations.

connect via neil.voishift@klyrr.com