Drift Monitoring for Voice Systems

How to stop silent degradation

A research-grounded whitepaper for teams building, buying, or operating voice systems for recorded calls, where performance can degrade quietly until trust and adoption collapse.

Version: v1.0

Publisher: VoiShift Research Series(neil.voishift@klyrr.com)

Abstract

Voice systems rarely fail in a single dramatic event. More often, they degrade quietly. A workflow that felt stable last month starts to feel slightly slower, slightly more interrupt-prone, slightly less accurate on domain terms, slightly more repetitive in recovery. Users do not file a precise bug report. They stop trusting the system, bypass it, or escalate earlier. That is silent degradation.

This whitepaper defines drift monitoring for voice systems as an operating discipline that detects and contains silent degradation before it becomes adoption decay. The discipline is grounded in the combined evidence base from the research sources, which collectively show that:

Pipeline & Holistic Evaluation

Real voice systems are a pipeline of components (speech detection, transcription, language processing, speech generation) that must behave as a single whole, and evaluation must go beyond simple accuracy metrics to include response timing, naturalness, and contextual understanding, because human perception determines whether users embrace or abandon voice technology. [3]

Turn-Taking & Overlap

Turn-taking and overlap behaviour is sensitive to information density and prosodic features, and human raters show clear preferences for overlap generation that accounts for both information density and suprasegmental cues, with results holding even with ASR errors. [2]

WER, SER & Security Metrics

Evaluation commonly relies on WER and SER, and SER is often neglected even though it can be 2–3 times higher than WER due to cumulative sentence errors; security evaluation adds metrics such as Attack Success Rate (ASR) to assess vulnerability and robustness. [4]

Endpoint Detection

Endpoint detection, frame overlap choices, and speech boundary extraction are foundational steps that influence downstream recognition and stability. [4]

Agentic Tool Monitoring

Tool integration in agentic systems requires monitoring and auditing, including logging all tool invocations, anomaly detection, and controlled rollouts via versioning and aliases, enabling traceability and debugging from proof-of-concept to production. [1]

Synthetic Voice & Provenance

Synthetic voice authenticity concerns and the limits of provenance safeguards affect trust conditions, with traceability mechanisms like C2PA and watermarks described as model-specific and limited in interoperability. [5]

Based on these findings, the paper offers a drift taxonomy, measurable drift signals, proof gates, a regression suite blueprint, and a monitoring programme that maps directly to how voice systems degrade in production.

Why Drift Is Different in Voice Systems

3.1 Drift is not only "accuracy got worse"

In voice systems, drift is often **felt before it can be measured** with a single metric. A team might notice more interruptions, more repeats, or more escalation requests before they notice an obvious jump in transcription errors.

This aligns with the view that voice evaluation must include human factors, not just accuracy. A voice agent architecture overview emphasises that testing demands holistic thinking beyond simple accuracy metrics and calls out response timing, naturalness, and contextual understanding as subjective elements that often determine whether users embrace or abandon voice technology. [3] That is a direct statement that "silent degradation" can happen in dimensions outside WER.

3.2 Voice drift is pipeline drift

Voice systems typically operate as a pipeline: speech-to-text, language model processing, and text-to-speech output, enabling conversational interaction. [6] A pipeline can drift even when each component still looks "within spec" individually. The problem is the interactions between components, especially around speech boundaries and turn-taking.

The security and privacy survey describes endpoint detection as extracting the beginning and conclusion of speech from a speech-containing signal, and it notes how time-domain and frequency-domain approaches have different failure trade-offs, including incorrect evaluation of air sounds and varied detection outcomes depending on volume calculations. [4] This is a concrete example of how small shifts in speech boundary behaviour can cascade into a system that feels less stable without a single "ASR model broke" moment.

3.3 Human conversation dynamics amplify drift

Turn-taking is not rigid in real speech. Work on information density and overlap explains that incremental systems process smaller chunks of user input to improve user experience and allow more flexibility in turn-taking. [2] The same work shows that humans prefer overlaps placed at low information density points and that these results hold even with ASR errors. [2] That matters because drift in overlap timing or turn boundaries changes perceived competence quickly.

In other words, a voice system can drift in ways that humans feel strongly even if the transcript remains "acceptable."

A Research-Grounded Definition of "Silent Degradation"

Silent degradation is a sustained, gradual change in the system's interaction behaviour or outcome quality that:

- does not trigger an immediate operational incident
- is not captured by a single obvious metric
- accumulates until users bypass, escalate, or stop trusting the workflow

This paper defines silent degradation using **five measurable symptom classes**, each supported by the evidence base:

1

Turn-Taking Degradation

Grounded in research showing overlap placement preferences and sensitivity to suprasegmental features and information density. [2]

2

Speech Boundary & Detection Degradation

Grounded in descriptions of endpoint detection and VAD as foundational, with varied outcomes and trade-offs [4] and the emphasis that effective systems adapt to changing environments rather than relying on rigid thresholds. [3]

3

Transcription & Sentence-Level Degradation

Grounded in WER and SER definitions and the note that SER can be 2–3 times higher than WER and is often neglected. [4]

4

Tooling & Orchestration Degradation

Grounded in guidance that production readiness includes monitoring capabilities, tracing, observability, logging tool invocations, anomaly monitoring, and controlled rollouts via versioning and aliases. [1]

5

Trust & Provenance Degradation

Grounded in the discussion that traceability mechanisms like C2PA and watermarking are model-specific with limited interoperability, and that detection and provenance gaps create trust fragility. [5]

Silent degradation is best treated as **drift across these five symptom classes**.

Why Teams Miss Drift Until Adoption Decays

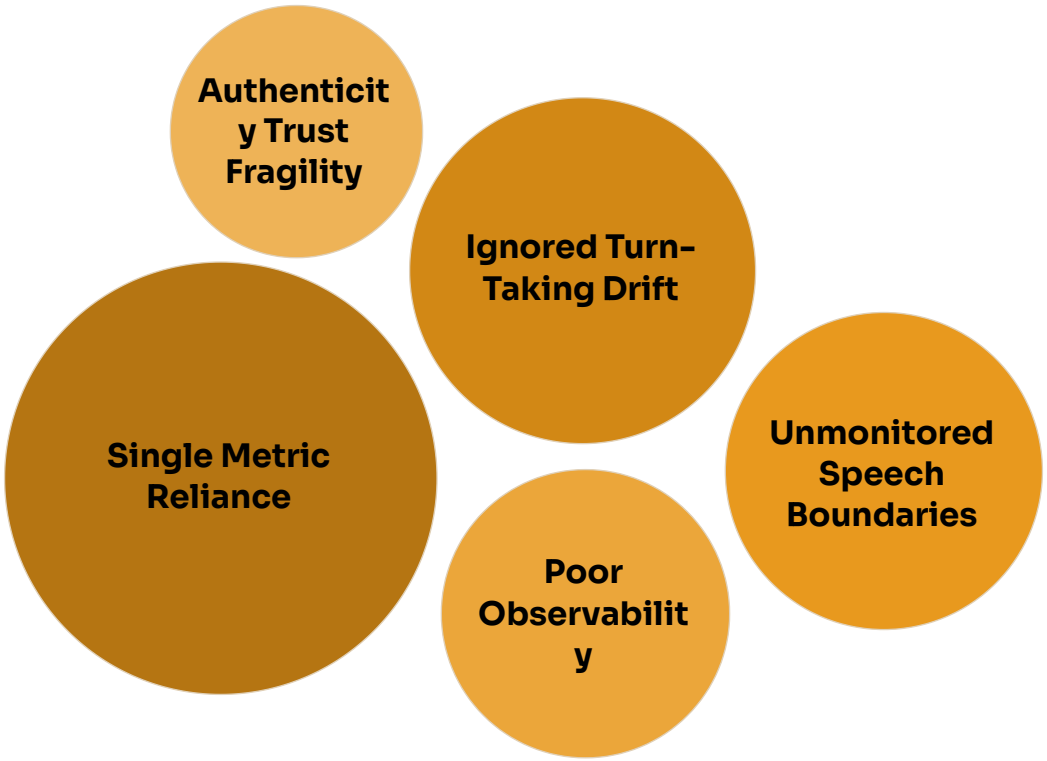

The five failure modes below explain why drift goes undetected until users have already changed their behaviour.

1 Over-reliance on one metric

WER is common. SER exists but is often neglected despite being higher because it accumulates errors at sentence level. [4] If a team tracks only WER, it can miss shifts in sentence-level stability that users experience as "the system keeps getting it wrong," especially in multi-step exchanges where one missing sentence changes the meaning.

2 Ignoring turn-taking drift

The overlap research shows that user preference is sensitive to features such as duration and pitch of overlapped segments and to information density peaks and troughs. [2] If a system's overlap and barge-in behaviour shifts slightly, it can feel less responsive or more interruptive, even if transcription accuracy does not move much.

3 Not treating speech boundary as a monitored control

Endpoint detection involves extracting speech boundaries and removing background noise in silent periods, with approaches that can lead to varied detection outcomes. [4] A voice system can slowly degrade if boundary detection shifts due to environment changes, device changes, or parameter updates, and most teams do not monitor that explicitly.

4 Weak observability and change control

The AWS agent systems guide calls out observability stack integration to enable traceability and debugging from proof-of-concept to production. [1] It also highlights monitoring and auditing practices such as logging tool invocations and monitoring for anomalies. [1] Without these, drift becomes "felt," not proven. That delays fixes and makes regressions repeat.

5 Trust fragility from authenticity uncertainty

The authenticity paper explains that provenance and traceability safeguards like C2PA integration and watermarking exist, but are model-specific and have limited interoperability, and not all providers implement them nor are they required to. [5] When trust conditions shift due to authenticity concerns, users can become less tolerant of minor errors. That changes the effective frustration threshold and makes small degradations more damaging.

Drift Taxonomy for Voice Systems

This taxonomy is designed to be operational. Each category is something you can monitor and test.

Category	Description	Key Signals / Source
6.1 Acoustic Environment	Changes in background noise, device mic quality, call routing, or room acoustics. Effective detection systems adapt as users move between environments, with thresholds set dynamically. [3]	Detection stability; environment context monitoring
6.2 Speech Boundary (VAD)	Endpoint detection approaches can produce varied detection outcomes and errors like mis-evaluation of air sounds. [4]	Increased cut-ins during pauses; late recognition; partial/fragmented utterances fed into ASR
6.3 Transcription (WER & SER)	SER often higher and neglected. [4] The correct utterance appears in the ASR N-best list only 41.7% when not ranked as 1-best; WER of 53.6 in CLASSiC corpus. [2]	Stable WER but rising SER on long turns; rising errors on specific term types; N-best instability
6.4 Turn-Taking & Overlap	Humans prefer overlaps at low information density points; significant preference ($p<0.001$) for combined suprasegmental and information density features. [2]	Overlaps at wrong points; backchannels too late/early; inconsistent barge-in handling
6.5 Orchestration & Tool	Monitoring and auditing, logging tool invocations, anomaly monitoring, rate limiting, tracing, and controlled rollouts via versioning and aliases. [1]	Changed tool behaviour; new retries/timeouts; inconsistent state between components
6.6 Security Posture	Attack Success Rate (ASR) is essential for assessing security and robustness. [4]	Increasing ASR trend; system becomes liability even if user-facing behaviour seems unchanged
6.7 Trust & Provenance	Safeguards like C2PA and SynthID are model-specific with limited interoperability and not universally implemented. [5]	Baseline trust shifts; same system behaviours become less acceptable as ecosystem changes

Core Principle: Drift Monitoring Must Cover Behaviour, Not Only Output

A voice system can degrade in perceived competence through timing and interaction behaviour even when the transcript remains acceptable. This is implied by:

- Evidence that turn-taking and overlap placement preferences are sensitive to evolving information density and prosodic features, with strong human preferences for combined-feature models. [2]
- The explicit statement that response timing, naturalness, and contextual understanding are subjective elements that determine adoption. [3]
- The reality that endpoint detection and speech boundary extraction are foundational steps with trade-offs that can lead to varied outcomes. [4]

So drift monitoring must include **behavioural signals** such as:

Interruption Success

Rate at which user barge-ins are correctly handled without cutting off meaningful speech

Cut-In Rate

Frequency of system cut-ins during user pauses — a direct signal of boundary detection drift

Repeat-Loop Frequency

How often users repeat themselves, indicating the system failed to process or respond correctly

Time-to-First-Response

Latency from end of user turn to system response onset — sensitive to orchestration and boundary drift

Repair Loop Count

Number of correction exchanges per session — a composite signal of transcription and turn-taking drift

These are not "nice-to-have." They are the **operational indicators of silent degradation**.

Measurement Framework: The Drift Scorecard

This paper proposes a Drift Scorecard with five dimensions. Each dimension maps to the taxonomy and has measurable indicators grounded in the sources.

Dimension A: Speech Boundary Stability

Indicators: endpoint detection failure markers (false starts, truncated turns); cut-in during pause rate; late turn closure rate.

Justification: Endpoint detection methods and their variability are described in the security survey. [4] The voice agent architecture overview emphasises dynamic adaptation to changing environments, implying boundary behaviour must remain stable across conditions. [3]

Dimension B: Transcription Stability

Indicators: WER trend; SER trend; N-best volatility proxy; domain-term error rate proxy (measured via targeted test sets).

Justification: WER and SER are defined, and SER being neglected is explicitly noted. [4] ASR ambiguity and error effects are discussed in the overlap paper, including N-best presence rate and WER figure for the dataset. [2]

Dimension C: Turn-Taking & Overlap Quality

Indicators: barge-in success rate; overlap placement appropriateness proxy (based on low vs high information density segments); backchannel timing consistency proxy.

Justification: The overlap and information density paper provides motivation for using information density to design and evaluate turn-taking modules and reports significant preference for combined feature models. [2]

Dimension D: Orchestration & Tool Integrity

Indicators: tool invocation audit completeness; anomaly detection alerts; rollback frequency; version alias drift correlation (behaviour changes after rollout).

Justification: The AWS guide specifies monitoring and auditing through logging tool invocations, anomaly monitoring, and rate limiting. [1] It also discusses tracing/observability and versioning/aliases for controlled rollouts. [1]

Dimension E: Trust & Security Posture

Indicators: authentication and authorisation failures; Attack Success Rate (ASR) trend where applicable; provenance coverage rate (what portion of outputs carry provenance metadata in supported systems).

Justification: ASR is described as essential for assessing security and robustness. [4] The authenticity paper explains provenance and traceability mechanisms and their limitations and the growth of third-party detection tools. [5]

Data Strategy for Drift Monitoring

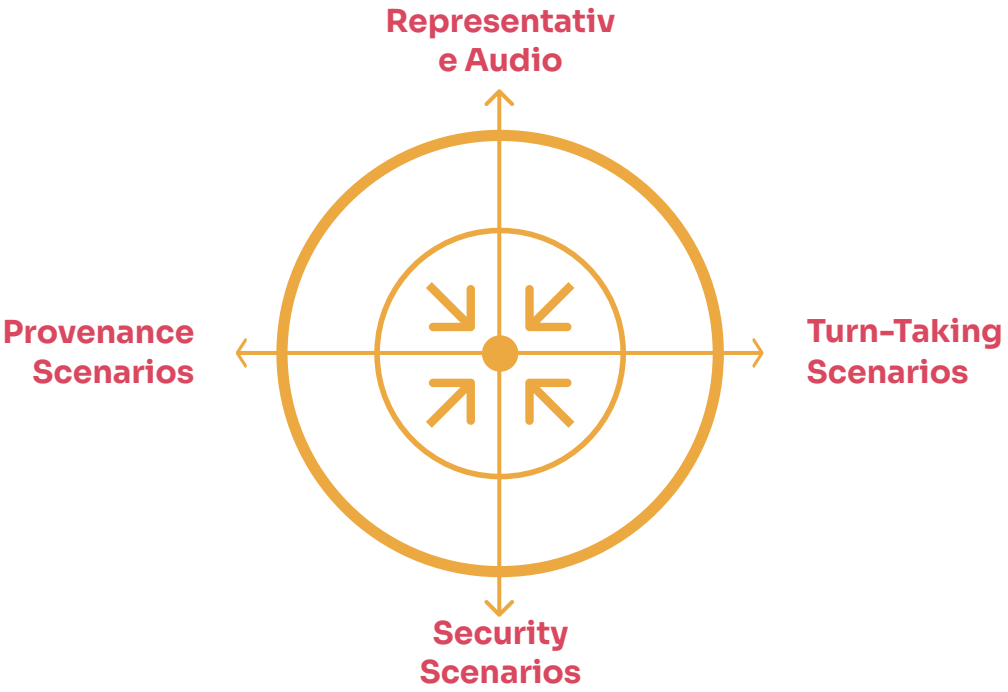

9.1 Why you need representative audio, not just logs

Voice drift is often environment-dependent. The voice agent architecture overview describes how moving from quiet to noisy environments transforms the acoustic landscape and how effective systems adapt without manual adjustments, suggesting that real-world diversity is a core challenge. [3] So you need: a representative audio set across environments and channels; a representative call set across workflows.

9.2 Why you need scenario libraries for turn-taking

Turn-taking is not captured well by random sampling alone because drift often appears in pauses, interruptions, corrections, and multi-speaker overlap. The overlap paper's emphasis on incremental processing and overlap phenomena implies you need scenario coverage that stresses those phenomena. [2]

9.3 Why you need a security scenario set

The security survey provides a taxonomy of security and privacy problems and emphasises ASR as a key metric for adversarial attempts. [4] That implies drift monitoring includes: periodic security evaluation runs; detection of changes in attack success indicators.

9.4 Why you need provenance scenarios

The authenticity paper notes that provenance and traceability safeguards are model-specific and not universally implemented. [5] This means drift monitoring should track whether provenance coverage is stable, especially if vendors or models change.

The Monitoring Loop: How to Detect Silent Degradation Early

This is a practical operating loop built from the monitoring and auditing principles described in the AWS guide [1] and the measurement needs described in the voice architecture and security sources. [3] [4]

1

Step 1: Establish Baselines

- Baseline WER and SER [4]
- Baseline boundary stability markers [4]
- Baseline turn-taking quality indicators [2]
- Baseline tool invocation patterns and audit completeness [1]
- Baseline security posture metrics including ASR where relevant [4]

2

Step 2: Instrument Continuously

- Log tool invocations, audit trails, anomaly detection [1]
- Log turn boundary events and interruption events
- Log transcription quality indicators and SER proxies
- Log provenance metadata presence where relevant [5]

3

Step 3: Run Scheduled Evaluation

Weekly: sample calls across environments and workflows; compute Drift Scorecard deltas.

Monthly: run scenario library regressions (turn-taking, boundary, security); review model and vendor changes using versioning/aliases. [1]

Quarterly: revalidate scenario library representativeness (new accents, new devices, new workflows). [3]

4

Step 4: Decide and Act

If drift is detected: isolate which dimension moved; roll back via versioning/aliases where possible [1]; patch boundary, orchestration, or model configuration; update regression suite so the drift cannot return.

Regression Suite Blueprint: How to Stop the Same Degradation from Coming Back

The goal of a regression suite is to prevent repeated drift caused by environment shifts, model updates, tool changes, and orchestration changes.

The overlap paper shows that human preferences are sensitive to overlap placement, information density, and suprasegmental features and that results hold even with ASR errors. [2] That implies regressions must include ASR error conditions, overlap conditions, and prosodic and density-sensitive conditions. The security survey's description of endpoint detection supports including boundary and segmentation tests and security regressions using ASR measures. [4] The AWS guide supports treating rollouts as controlled and observable through tracing and versioning/aliases. [1]

Suite Component	What to Test	Source Grounding
1. Boundary Tests	Pauses, fillers, background noise; endpoint detection robustness checks; VAD behaviour across environments	[4] [3]
2. Turn-Taking Tests	Overlap placement tests based on low vs high information density segments; interruption and barge-in handling tests	[2]
3. Transcription Tests	WER and SER tests; long-turn SER stress tests since SER can be higher and neglected; N-best ambiguity scenarios consistent with ASR realities	[4] [2]
4. Orchestration Tests	Tool invocation audit completeness; timeout and retry behaviour consistency; anomaly detection triggers on unusual tool use patterns	[1]
5. Security Tests	Representative adversarial tests where relevant; ASR trend tracking; rate limiting and abuse prevention hooks	[4] [1]
6. Provenance Tests	Verify provenance metadata presence and stability when using supported systems like C2PA; verify detection classifier behaviour changes do not create new false positives or negatives	[5]

Drift in "Human Feel" and Why It Must Be Monitored

The voice agent architecture overview explicitly states that response timing, naturalness, and contextual understanding are subjective elements that determine whether users embrace or abandon voice technology. [3] That means "human feel" is not a brand topic only. It is a **monitored quality attribute**.

The overlap research adds a mechanism-level explanation: humans are sensitive to peaks and troughs in evolving information density and prefer overlaps at low information density points, and these results hold even with ASR errors. [2] That means "human feel" can degrade when:

- overlap comes at the wrong time
- the system gives feedback at high-density moments
- the system's responsiveness becomes poorly aligned to conversational load

Monitoring "human feel" in this whitepaper is not framed as aesthetic scoring. It is:

Turn-Taking Quality Changes

Detecting shifts in overlap placement, barge-in handling, and backchannel timing relative to information density

Timing & Interruption Success

Detecting changes in time-to-first-response and whether interruptions are handled correctly

Repeat & Repair Loops

Detecting changes in repeat loops and repair loops — observable markers of perceived system failure

Restatement Frequency

Detecting changes in how often the user has to restate information — a direct signal of comprehension drift

These are **observable markers**, not subjective impressions.

Worked Example: Silent Degradation in a Medium-Risk Workflow

Status updates and exception handling. This example is intentionally medium-risk because it is common in support, sales, and ops and can trigger escalation without directly moving money.

Baseline Behaviour (Week 1)

- Turn boundaries stable
- Quick response start
- Few repeat loops
- Clear handling of uncertain state

Drift Onset (Week 6)

Users begin:

- repeating "is it arriving today or not"
- interrupting more often
- requesting a human sooner

The system metrics show:

- no obvious large jump in WER
- subtle increase in boundary errors (more cut-ins during pauses)
- subtle increase in tool latency dead air
- slightly higher SER on longer turns

Why This Pattern Is Plausible Within the Evidence Base

- SER can be neglected despite being 2–3 times higher than WER due to cumulative errors. [4] That can hide sentence-level drift.
- Boundary detection and endpoint extraction are foundational and can produce varied outcomes depending on approach and conditions. [4]
- Users are sensitive to overlap and feedback timing; human preference shifts based on information density and suprasegmental features. [2] Small turn behaviour shifts can be felt quickly.
- Voice evaluation must include timing and human perception factors because they determine abandonment. [3]

Drift Response Using the Scorecard

Dimension A boundary stability worsened: increase in cut-in rate during pauses (endpoint detection drift). [4]

Dimension D orchestration integrity worsened: increased dead air during tool calls (orchestration timing drift). [1]

Dimension B transcription stability worsened slightly: SER increased even if WER did not move much. [4]

Intervention Plan

- Patch endpoint detection tuning or environment adaptation path. [4]
- Reduce orchestration dead air via monitored tool invocation latency patterns, with audit logs and anomaly detection. [1]
- Add SER-focused regression tests to prevent recurrence. [4]

This example shows what silent degradation looks like operationally: **users change behaviour first**, and basic WER-only dashboards miss it.

Security Drift and Trust Collapse

Silent degradation is not always "the system got worse at listening." Sometimes it is "the ecosystem changed and trust collapsed."

Provenance & Authenticity Mechanisms

The authenticity paper describes provenance and traceability mechanisms, including C2PA and watermarking (SynthID), but emphasises that safeguards are model-specific and have limited interoperability and are not universally implemented nor required. [5] It also discusses the rise of third-party synthetic voice detection tools, including "synchronous" detection capabilities or "liveness testing" for contact centres in banking, finance, insurance, healthcare, or retail. [5] This shows that security and authenticity concerns can change adoption conditions.

Attack Success Rate as a Core Metric

The security survey defines Attack Success Rate (ASR) as measuring the effectiveness of attacks by the proportion of successful adversarial attempts to manipulate or deceive speech recognition and calls it essential for assessing security and robustness. [4] This supports a core operating requirement: drift monitoring must include security posture drift, because an increase in attack success indicators can cause immediate trust loss even if the call experience still "works."

Operational Requirements

- Track ASR trend where relevant. [4]
- Log tool invocations and monitor anomalies. [1]
- Apply rate limiting to prevent abuse via excessive tool calls. [1]
- Maintain secure communication and data minimisation during tool integration. [1]

> Security drift is silent until it is catastrophic. Monitoring exists to detect it earlier.

Vendor and Model Drift: Why Controlled Rollouts Are Part of Drift Monitoring

Drift often enters through changes: new model versions, new tool endpoints, new orchestration logic, new vendor behaviour.

Observability and Versioning as Drift Containment

The AWS guide explicitly describes tracing and observability to track step-by-step reasoning, and versioning and aliases to create multiple versions and deploy them through controlled rollouts. [1] It also describes observability stack integration (CloudWatch, AgentOps, LangFuse) for traceability and debugging from proof-of-concept to production. [1]

These are not "nice platform features." They are **drift containment mechanisms**.

Programme Requirements

A drift programme should require:

- any behaviour change ships behind an alias
- every alias has baseline scorecard metrics
- every alias has regression suite pass results
- any anomaly triggers rollback, not debate

This keeps silent degradation from becoming a multi-month investigation with no reproduction path.

Drift Entry Points

- New Model Versions
- New Tool Endpoints
- New Orchestration Logic
- New Vendor Behaviour

Implementation Checklist: Minimum Viable Drift Monitoring Programme

This checklist is a strict minimum derived from the sources above.

Instrumentation

- Log all tool invocations and maintain audit trails. [1]
- Monitor for anomalies in tool usage patterns. [1]
- Implement rate limiting where applicable. [1]
- Capture turn boundary events and interruption events. [4]
- Track WER and SER, including SER explicitly. [4]

Evaluation

- Maintain a scenario library for overlap and turn-taking stress tests. [2]
- Include boundary detection tests under varying environments. [3] [4]
- Include sentence-level tests and long-turn SER tests. [4]
- Include security evaluation and ASR tracking where relevant. [4]

Release Discipline

- Use versioning and aliases for controlled rollouts. [1]
- Use tracing and observability to reproduce "felt wrong" issues. [1]

Trust and Provenance

- Track provenance and traceability coverage where relevant, acknowledging model-specific and limited interoperability constraints. [5]
- Treat authenticity concerns as a monitored trust axis, not an external PR issue. [5]

If you do not have these, **silent degradation will be discovered only after users have already decided to bypass.**

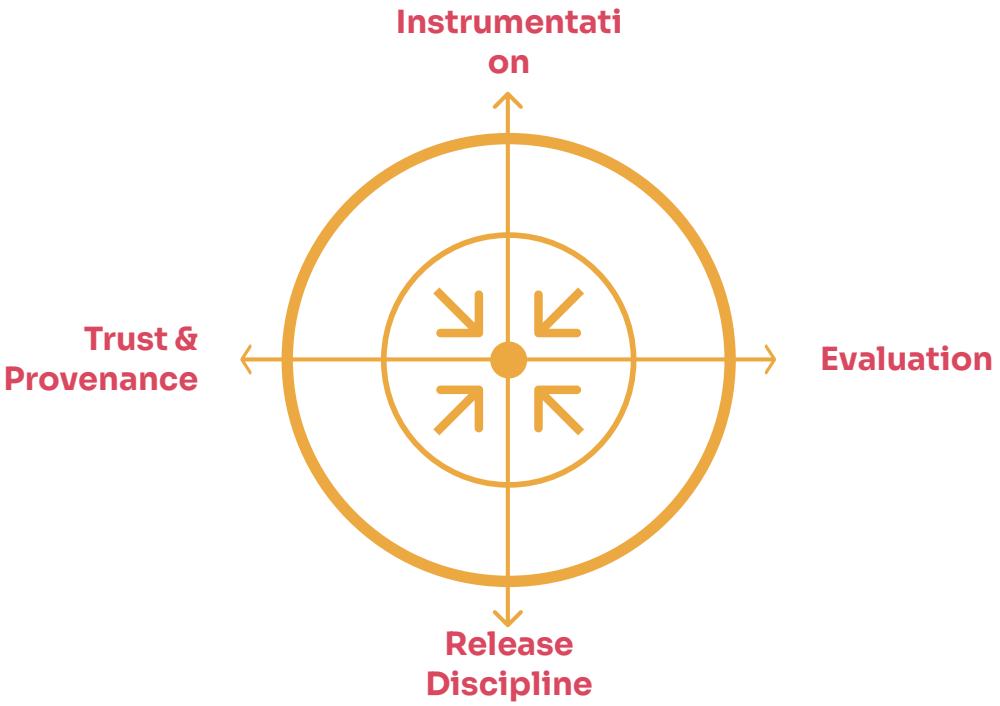

Conclusion

Silent degradation is predictable in voice systems because voice is a pipeline operating under variable conditions, and human perception is a primary success criterion. The evidence base supports that:

Holistic Testing Beyond Accuracy

Holistic testing beyond accuracy metrics is needed because response timing, naturalness, and contextual understanding often determine whether users embrace or abandon voice technology. [3]

Turn-Taking Design & Evaluation

Turn-taking and overlap behaviour can be designed and evaluated using information density and suprasegmental features, with human raters showing significant preference for combined-feature overlap models and preference for overlap at low information density points, even with ASR errors. [2]

SER Alongside WER

Evaluation should include SER alongside WER because SER can be 2–3 times higher and is often neglected, and security evaluation should include ASR as an essential robustness metric. [4]

Monitoring & Auditing Practices

Monitoring and auditing practices such as logging tool invocations, anomaly monitoring, rate limiting, and controlled rollouts via versioning and aliases enable traceability and drift containment. [1]

Provenance & Trust Conditions

Provenance and traceability safeguards exist but are model-specific and have limited interoperability, and the rise of detection tools reflects changing trust conditions that can reduce user tolerance for errors. [5]

> Drift monitoring for voice systems is therefore not a single dashboard. It is an operating discipline with a scorecard, regression suite, and controlled release process that detects and contains silent degradation before it becomes adoption collapse.

Bibliography

The following sources form the evidence base for this whitepaper. All citations in the document refer to these works.

Ref	Full Title & Authors	Publisher / Journal
[1]	Sempf, Aaron; Samuel, Joshua. *Agentic AI frameworks, protocols, and tools on AWS.* July 2025.	AWS Prescriptive Guidance, Amazon Web Services (AWS)
[2]	Dethlefs, Nina; Hastie, Helen; Cuayáhuitl, Heriberto; Yu, Yanchao; Rieser, Verena; Lemon, Oliver. *Information Density and Overlap in Spoken Dialogue.* Preprint submitted October 30, 2015.	*Journal of Computer Speech and Language*
[3]	Paul, Varghese. *Voice AI Agents: The Technical Backbone of Modern Speech Interfaces.* Vol 5, Issue 7, 2025.	Sarcouncil Journal of Multidisciplinary (SARC Publisher)
[4]	Li, Jingjin; Chen, Chao; Rahimi Azghadi, Mostafa; Ghodosi, Hossein; Pan, Lei; Zhang, Jun. *Security and privacy problems in voice assistant applications: A survey.* 2023.	*Computers & Security* (Elsevier)
[5]	Burgess, Jean; Carlon, Dominique; Doyuran, Elif Buse; et al. *Voice AI and authenticity: Current issues and emerging challenges.* Working Paper 012. August 2025.	ADM+S Working Paper Series
[6]	*The Rise of Voice AI: Transforming Human-Machine Conversations.* Special Report. April 2025. Source cited in report: Cartesia.	Special Report
[7]	Zhang, Zheng. *Comprehension Without Competence: Architectural Limits of LLMs in Symbolic Computation and Reasoning.* 11/2025.	Transactions on Machine Learning Research (TMLR)

This whitepaper is published by **VoiShift Research Series**, Version v1.0, dated 2026-02-22. All claims are grounded in the sources listed above. No invented claims, examples, numbers, or names have been introduced.