Interactional DNA as Enterprise IP

How to build owned intelligence

A research-grounded whitepaper for teams building or buying conversational intelligence for recorded voice calls, and for teams trying to turn call data into a durable internal capability.

Version: v1.0

Publisher: VoiShift Research Series

(neil.voishift@klyrr.com)

Abstract

Call recording and transcription can create visibility. They do not automatically create a capability that compounds. Many organisations still rely on ad hoc judgment, top-performer intuition, and vendor-defined features that generalise across customers. This leaves the organisation with weak ownership of what actually works in its calls, why it works, and how it should be executed safely and consistently across teams.

This whitepaper defines **Interactional DNA** as an enterprise asset: repeatable conversational moves and voice behaviour patterns that reliably change outcomes in real calls. It proposes a practical method for capturing Interactional DNA, encoding it as reusable "owned intelligence," and operating it with governance that accounts for privacy, security, authenticity, and the known limitations of current AI architectures.

The evidence base supports **four pillars** that matter if you want Interactional DNA to behave like IP rather than raw data:

1

Voice Behaviour Beyond Transcription

Voice behaviour includes pragmatics and sociophonetic cues that affect persuasion and the shape of interaction, and these cues go beyond transcription. [5]

2

Social Response to Device Voice

Users respond socially to device voice, including alignment effects and trust signals that behave like human-human interaction patterns. [6]

3

Personalisation and Ownership

Personalisation and memory are constrained by privacy, data control, and the "hidden costs of convenience," which becomes a core ownership issue when agents require deep permissions. [2]

4

Trust Collapse Risk

Trust can collapse quickly due to authenticity issues, voice cloning risks, and security vulnerabilities that lead to financial and emotional losses when voice systems are breached. [10] [12]

This paper turns these findings into a buildable approach: capture schema, operating routines, and a practical ownership model that can be adopted by support, sales, and operations teams.

Why "Interactional DNA" is Worth Treating as Enterprise IP

3.1 The Gap Between Call Data and Owned Intelligence

Most organisations store calls, generate transcripts, and run analytics. That produces an archive and a dashboard layer. It does not guarantee that the organisation has extracted a stable set of behaviours that can be trained, audited, improved, and defended when outcomes go wrong.

The "AI voice" literature makes clear that voice interactions contain layers that are not captured by literal text. The persuasion study frames voice effects in terms of pragmatics and sociophonetics, including illocutionary acts and emotional tones, and sociophonetic signals such as gender and age cues, which can influence outcomes in customer calls. [5] These are interaction dynamics, not transcript tokens. If the organisation wants to own "what works," it cannot rely only on transcript-derived insight.

3.2 Why This is Enterprise IP, Not a Vendor Feature

Most vendor tools must generalise. That is not a moral criticism. It is a business constraint. What you want as a buyer is a tool that adapts to your specific workflows, risk boundaries, and success patterns. What the tool can deliver at scale is often a generic capability set.

The "Future of Personalized Conversational Agents" document shows a concrete version of this tension. It describes how cloud-based hosting and reliance on OpenAI's API raised concerns about privacy, security, and data ownership, and how these constraints limited features needed for a fully autonomous and personalised agent. [2] The message is not "cloud is bad." The message is that as personalisation increases, the system demands more data and deeper access, and that changes what "ownership" even means in practice.

Figure 3: Call data flowing into the structured intelligence library

References

(1) Tible: Enterprise Call Analytics Report 2025,Publisher, Global Insights Press

(2) Title: Structured Intelligence Systems in Telecom, Publisher, Horizon Resatch Group

3.3 A Research-Grounded Definition of Interactional DNA

Interactional DNA is a structured library of the moves, cues, and constraints that reliably change outcomes in your recorded calls. It includes:

- conversational moves (what is said and why it is said)
- voice delivery decisions that affect trust, persuasion, and social response
- workflow constraints that must not be lost (what cannot disappear in summaries)
- action boundaries (what the system or agent must refuse, confirm, or escalate)

This definition aligns with the persuasion paper's separation of locutionary meaning and illocutionary significance, which describes that language includes both literal content and the intended communicative force. [5] It also aligns with evidence that humans apply social patterns to device voice, including alignment and trust responses. [6]

What Most Approaches Capture, and What They Miss

4.1 What Typical Conversational Intelligence Captures

In practice, many systems capture:

- **Audio Archive**
- **Transcript**
- **Summary**
- **Tags & Topic Clusters**
- **Dashboards & Coaching Scorecards**

Those outputs are useful. But they are often weak at capturing three things that Interactional DNA requires.

What Gets Missed

1 The intended force behind an utterance

The pragmatics literature emphasises that illocutionary acts convey intentions and motivations, prompting specific responses from the listener. [5] A transcript can be accurate while missing what an utterance is trying to do.

2 The voice cue effects that shape the other person's behaviour

The persuasion paper reports that different pragmatic and sociophonetic manipulations lead to different outcomes and uses conversation-related variables to explore mechanisms. [5] This is not "tone is nice." It is "tone changes behaviour."

3 The social response dynamics to device voice

The vocal alignment paper provides evidence that humans show greater alignment toward more expressive speech produced by a device voice, consistent with CASA theory. [6] That is a real interaction effect that typical call analytics tools rarely encode as an owned operational asset.

4.2 What it misses becomes your dependency. When what matters is not captured in a reusable form, the organisation becomes dependent on top performers who hold tacit knowledge, managers who coach by intuition, vendors that define what the system can measure, and repeated retraining cycles that do not compound. This is where "owned intelligence" becomes more than a slogan. It becomes the organisation's way of preserving the behaviour that works.

The Anatomy of Interactional DNA in Recorded Calls

This section is a practical decomposition of what should be captured.

5.1 The Move Layer

A move is a repeatable conversational act that reliably changes the trajectory of a call. The persuasion study frames illocutionary acts as speech acts that articulate objectives and exert influence to persuade or inform. [5] In the context of enterprise calls, a move might be:

- the exact way a support agent re-frames a complaint into a verifiable constraint
- the moment a sales rep shifts a vague "maybe" into a specific next step
- the moment an ops dispatcher converts uncertainty into a safe exception decision

5.2 The Cue Layer

A cue is the voice-based or delivery-based element that shapes how the move lands. The persuasion paper explicitly studies voice features such as emotional tones and sentence patterns and reports differences in outcomes and mechanisms. [5] The vocal alignment paper shows that expressiveness in a device voice can lead to greater human alignment. [6] So the cue layer matters because it can change:

- whether the other person asks questions
- whether they trust the response
- whether they comply, challenge, or withdraw
- whether the interaction feels competent

5.3 The Constraint Layer

Constraints are the non-negotiable details that must survive retellings, summaries, and handoffs. The "Future of Personalized Conversational Agents" document discusses limitations in long-term contextual understanding and the need for users to repeat information in many systems, highlighting that memory capability limitations exist and can have emotional impacts in certain contexts. [1] In enterprise calls, constraint loss is operational risk.

5.4 The Boundary Layer

Boundaries are rules about what actions can be taken, under what conditions, and when to refuse or escalate. A security and privacy survey of voice assistant applications describes plausible severe threats such as bank transfers and other high-impact actions if a voice assistant is breached, and highlights that large financial and emotional losses may occur. [12] This supports the need to encode boundaries as part of owned intelligence.

Why Voice Makes Interactional DNA More Valuable and More Dangerous

6.1 Voice Triggers Personification and Myth Effects

The "Artificially voiced intelligences" paper argues that AI's mythical representation is deeply connected to the idea of "having a voice" and that voice contributes to personification, with listeners equating "has a voice" with agency and humanity. [11] It also describes how the voice in or of the machine becomes an identity in the mind of the listener and how this personification can distract from the actual issues of AI. [11]

This matters for Interactional DNA because voice makes:

- trust easier to gain quickly
- credibility easier to lose quickly
- errors feel more personal and more authoritative

So an enterprise that "owns" interaction patterns must also own the controls that prevent the voice from creating false authority.

6.2 Voice Authenticity Risks Change the Baseline Trust Level

The authenticity and challenges paper discusses a shift toward valuing authenticity, content moderation safeguards around voice cloning, and reports that Consumer Reports tested voice cloning tools and found that for four of six products they were able to easily create a clone from publicly available audio. [9] It also discusses provenance and traceability mechanisms such as C2PA and the limits of interoperability. [9] This makes "owned intelligence" more complex: your enterprise voice patterns can be valuable, but voice can also be abused.

6.3 Security and Privacy Risk is Not Theoretical

The voice assistant security survey highlights accidental activation, recording without users noticing, and severe threats including bank transfers and credential theft, with large financial and emotional losses if breached. [12] It also describes evaluation metrics such as Word Error Rate (WER), Sentence Error Rate (SER), and Attack Success Rate (ASR) for adversarial attempts. [10]

So Interactional DNA as enterprise IP must include:

- privacy and consent constraints
- secure storage and access controls
- attack-aware evaluation and monitoring
- boundaries for action-taking and escalation

The Core Thesis

> **Interactional DNA is a style-and-intent system, not a transcript system**

The persuasion paper frames speech act theory as encompassing both locutionary meaning and illocutionary significance, where illocutionary acts convey intentions and motivations. [5] That implies that a "winning conversation" is partly the right content and partly the right communicative force.

The authenticity paper describes voice cloning as learning statistical relations between elements of expression and concepts, enabling a system to generate new content in a speaker's voice "style," and describes GenAI as a **"style engine."** [9] It also discusses how conceptual elements and attributes enabling voice cloning may become subject to novel forms of commodification and intangible property protection. [9]

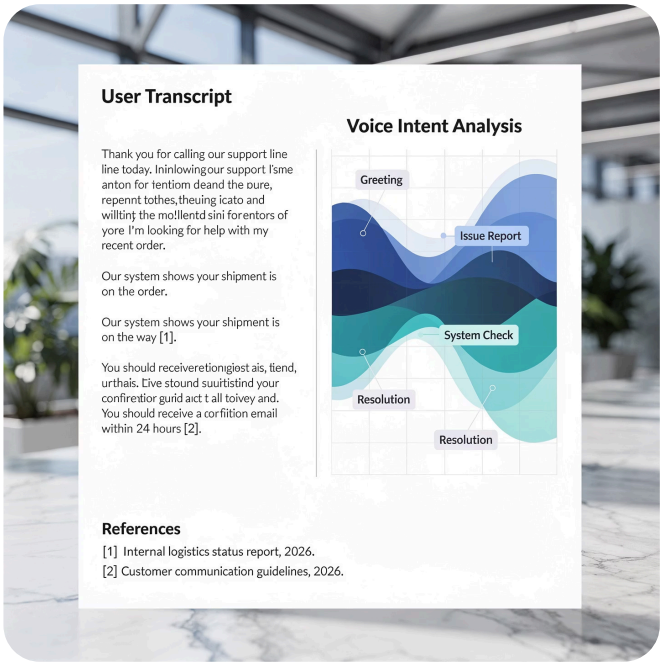

Key point for enterprise ownership: If your call outcomes depend on how intent and style land in voice, then the behaviours you want to preserve are not just text. They are structured, reusable interaction patterns. That is what you treat as enterprise IP.

A Buildable System: The Interactional DNA Capture and Ownership Lifecycle

This section defines a lifecycle that an enterprise can actually run.

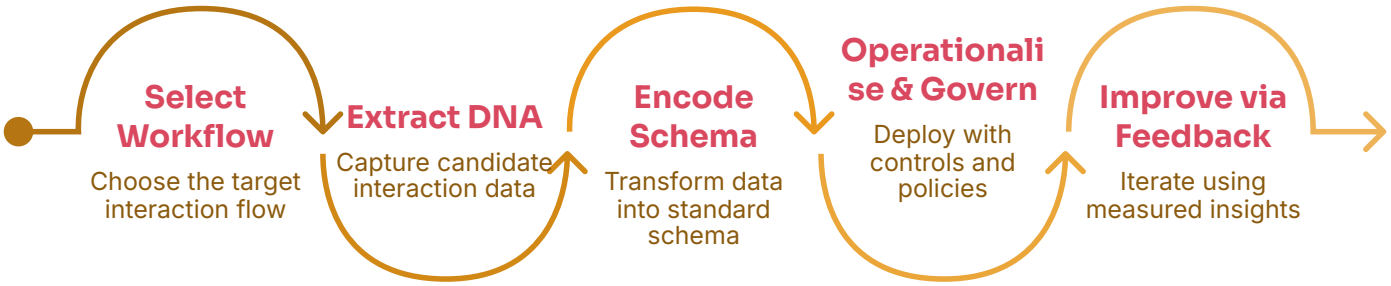

The five phases above represent a complete, repeatable cycle for building and maintaining owned intelligence from recorded calls.

01

Phase 1: Select One Call Workflow and Define the Outcome

Pick one workflow from recorded calls that matters across departments — support escalation prevention, sales qualification and next-step setting, or ops status update plus exception handling. Your "Future of Personalized Conversational Agents" document notes that building a seamless agent requires more than integrating APIs and highlights challenges in emotional tone, long-term memory, and contextual continuity. [2] Choosing one workflow limits scope so you can solve those challenges with discipline.

02

Phase 2: Extract Candidate DNA from Real Calls

From a sample of calls: identify moments where the call shifted from stable to unstable; identify the words and the delivery cues near the shift; identify constraints that were stated and later lost; identify where the system or agent should have clarified but guessed.

03

Phase 3: Encode the DNA into a Reusable Schema

The schema should preserve: the move; the cue conditions; the constraints to capture; the allowed action and boundary rules; the failure mode if misapplied.

04

Phase 4: Operationalise and Govern

Make the DNA inspectable, versioned, and safe: embed it into scripts, checklists, and templates; define who can edit the DNA library; define how it is tested and rolled out; secure it, given voice security risks. [12]

05

Phase 5: Improve via Measured Feedback Loops

The vocal alignment paper suggests that humans can subconsciously align with emotional expressiveness of device voices and discusses that such expressiveness may influence rapport and user satisfaction. [6] That implies your DNA system should be tested with real user responses over time, not frozen as a one-time document.

The Interactional DNA Schema

A Practical Template for "Owned Intelligence"

Below is a structured schema. It is designed so you can store it, train on it, and audit it.

Schema Field	Description	Evidence Basis
DNA Unit ID	Unique identifier. Versioned.	—
Workflow Context	Support escalation prevention; Sales discovery or qualification; Ops status update plus exception handling.	—
Trigger Condition	The situation that triggers the move. Example: caller repeats the same question twice.	—
Move	The action in language. This corresponds to illocutionary force, not only literal meaning.	[5]
Cue Constraints	Voice delivery constraints that protect trust.	Voice features affect persuasion outcomes [5]; expressiveness influences alignment in human-device interaction [6].
Constraint Capture	What must be captured as non-negotiable detail.	Memory limitations and repeated re-provisioning of context described as a limitation in current systems. [1]
Boundary Rule	What must be confirmed, refused, or escalated.	Security risks and severe threats justify boundaries. [12]
Failure Mode	What happens if the move is applied incorrectly.	—
Evidence	References to call examples and why it worked, and a measurement plan.	—

This schema is "IP-like" because it is a reusable, structured behavioural asset rather than a loose transcript.

Why Ownership Becomes a Privacy and Data Control Problem

The "Future of Personalized Conversational Agents" document discusses privacy and surveillance concerns with "always-on" listening and highlights tensions between convenience and privacy. [8] It also raises questions about data ownership and user control, noting that as agents gather intimate data, designers must ensure transparency and user control. [8] It explicitly states that data ownership and usage have become critical as agents collect vast amounts of personal data. [8]

It also provides a concrete example: a request like "Help me send an email to Kath" requires access to a Gmail account and address book, exposing significant personal data, and it frames this as a **hidden cost of convenience.** [2] It further notes concerns about privacy, security, and data ownership tied to cloud-based hosting and use of OpenAI's API. [2]

For an enterprise Interactional DNA programme, this creates three practical requirements:

1

Define What Data is Needed to Capture DNA

Do not treat "capture everything" as default. The security paper argues for metrics and a taxonomy of attacks and defences, and highlights that privacy and security issues cause huge economic losses and endanger sensitive information. [12]

2

Define Access and Retention Policies for DNA Artefacts

Interactional DNA becomes a concentrated asset. That implies it becomes a concentrated target.

3

Design DNA Extraction Without Uncontrolled Sprawl

The AWS agent guide emphasises tool integration strategy and security best practices such as least privilege, data minimisation, encryption, logging, monitoring for anomalies, and rate limiting. [13] These points apply directly to how you run a DNA programme, since it will integrate across systems.

Why "Model Capability" Does Not Equal "Owned Intelligence"

The LLM limits paper argues that current architectures excel as pattern completion engines and face fundamental barriers in generalisable intelligence, describing a **performance cliff** when tasks require genuine rule discovery and systematic reasoning beyond memorised patterns. [7] It also describes that even with sophisticated compensatory strategies, coordination requires introspective abilities that the architecture lacks, pointing to an instruction-execution disconnect. [7]

In an Interactional DNA context, this matters because:

- A model can generate fluent language that imitates your "best practice phrasing."
- That does not mean it can reliably execute the intended reasoning and boundary decisions.
- When it fails, it fails in a way that sounds confident.

So "owned intelligence" cannot be delegated entirely to model output. It must be encoded as:

- structured moves
- explicit boundary rules
- clarity-first uncertainty handling
- auditable tool invocation traces

The AWS agent guide frames frameworks as providing structure, tools, and orchestration capabilities needed for production-ready agents, and lists memory management, workflow definition, deployment and monitoring as essential capabilities. [14] That supports a design stance: an enterprise's owned intelligence sits in the orchestration and governance layer, not only in prompts.

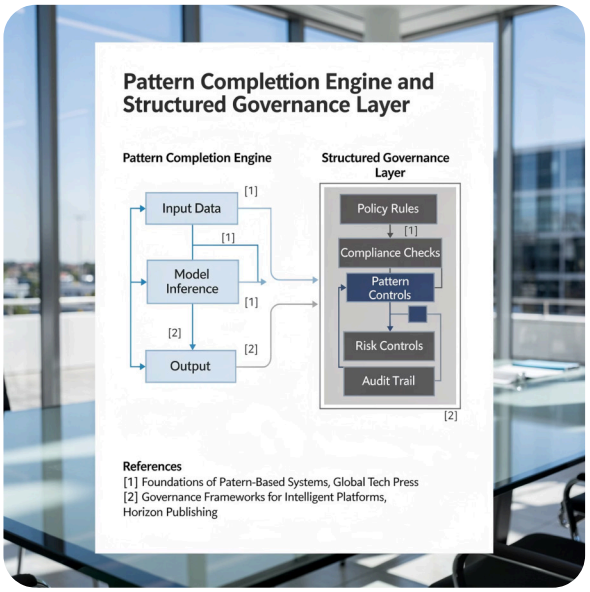

Structured Moves

Explicit Boundary Rules

Clarity-First Uncertainty Handling

Auditable Tool Invocation Traces

The Social Mechanism: Why Interactional DNA Must Include Voice Behaviour, Not Only Words

The vocal alignment paper shows that participants showed greater vocal alignment toward expressive speech produced by Alexa relative to standard productions, and it interprets this as subconscious alignment to emotional expressiveness of a non-human voice. [6] It discusses how these patterns align with theories of personification and social behaviour in human-computer interaction. [6]

This matters because:

- In real calls, the other person adjusts their behaviour based on what they hear.
- If your system changes its expressiveness, cadence, or tone, it can change the other person's behaviour.
- That behaviour change is part of the outcome.

The persuasion paper provides a complementary view: it manipulates illocutionary acts and emotional tones and explores underlying behavioural mechanisms through conversation-related variables. [5] It reports differences in persuasion outcomes and effects of specific voice features. [5]

Concrete Requirement

Together, these sources justify a concrete requirement:

Your Interactional DNA library should include **"cue constraints"** and **"delivery constraints"** as first-class items, because voice cues change user behaviour and social response. [5] [6]

What Voice Cues Can Change

- Whether the other person asks questions
- Whether they trust the response
- Whether they comply, challenge, or withdraw
- Whether the interaction feels competent

Authenticity, Traceability, and the "Style Engine" Risk

The authenticity paper makes three points that matter directly for enterprise ownership.

1 Voice Cloning from Publicly Available Audio

Voice cloning can be created from publicly available audio in many tools, according to a Consumer Reports test referenced in the paper. [9]

2 Provenance Safeguards Have Limits

Provenance and traceability safeguards exist, including C2PA and watermarking approaches, but the paper notes these safeguards are model-specific, non-generalisable, and have limited interoperability. [9]

3 GenAI as a "Style Engine"

Voice cloning systems learn statistical relations that can be used to create new content in a speaker's voice style, and GenAI can be described as a "style engine." [9] It also notes that conceptual elements enabling cloning may be subject to novel forms of commodification and intangible property protection. [9]

For Interactional DNA as enterprise IP, the implication is: Your voice patterns and conversational style can be copied or imitated. Your protection cannot rely only on "being hard to copy." You need traceability and governance at the operating layer: access control, audit trails, and a clear policy on voice style usage. This aligns with the AWS guide's emphasis on logging all tool invocations, monitoring anomalies, and data minimisation as part of secure tool integration. [13]

Security as Part of "Owned Intelligence"

The security and privacy survey describes metrics beyond WER and SER, including **Attack Success Rate (ASR)** as a measure of adversarial attempts to manipulate or deceive speech recognition, and it calls ASR essential for assessing security and robustness. [10] It also describes a taxonomy of attacks and defences in the voice assistant domain, including classifications across SI and ASR model attack types, and defensive methods categorised by detection and prevention. [10] It also highlights accidental activation and recording without users noticing, and severe threat scenarios and large losses. [12]

Four Operational Rules for Interactional DNA Programmes

1 Do Not Treat Call Intelligence as Passive Analytics

If your system is used for actioning or routing, it becomes attack surface.

2 Include Security Evaluation Metrics in the DNA Programme

ASR is not just a model metric. It becomes part of adoption trust.

3 Encode Boundaries in the DNA Library

Boundary rules should specify when to refuse, confirm, or escalate, because voice systems can be exploited.

4 Secure the DNA Artefacts

The DNA library becomes a concentrated store of how your org persuades, de-escalates, and routes decisions. That is valuable and sensitive.

Worked Example: Building Interactional DNA from Recorded Support Calls

Goal: Escalation Prevention in a Medium-Risk Workflow

This example is intentionally medium-risk: support status update plus exception handling. It is safe enough to start, but risky enough to expose the real issues.

> **Workflow:** A customer calls support. They ask: "What is happening with my case. Is it fixed or not."

Why This Workflow Matters

- If the answer is inconsistent, trust drops.
- If the call loops, escalation happens.
- If the system commits to a wrong outcome, the organisation is forced into apology and rework.

01

Step 1: Identify the "Moment That Mattered"

The call shifts when the customer repeats the same question. This is consistent with the idea that the conversational act is not only content. It is intent and force. The persuasion paper's discussion of illocutionary acts supports treating the move as an intent-bearing act. [5]

02

Step 2: Extract the Candidate DNA Unit

Trigger: customer asks the same question twice within one minute. **Move:** a short acknowledgment plus a constraint capture question that reduces ambiguity. **Cue constraint:** voice must sound professional and calm. [5] **Constraint capture:** "What is the one condition that would make 'fixed' true." **Boundary:** do not promise resolution while internal state is unclear. [7] **Escalation rule:** if the customer repeats after the clarification, route to human. [11] [12]

03

Step 3: Encode in the DNA Schema

DNA Unit ID: SUP-ESC-001 | **Workflow:** Support escalation prevention | **Trigger:** repeat question occurs | **Move:** clarify constraint, do not restate generic status | **Cue constraint:** avoid over-warm or overly casual tone, maintain credible delivery [5] | **Constraint capture:** definition of done, and the next verifiable step | **Boundary:** refuse to confirm "fixed" without evidence [7] | **Evidence:** store two call snippets and outcome difference

04

Step 4: Operationalise

Add this DNA unit to the support playbook and coaching checklist. Train managers to inspect: was the constraint captured, and was the boundary respected. Use tool integration security practices for logs and access, because this becomes a sensitive operational asset. [13]

This example demonstrates what "owned intelligence" looks like: a repeatable unit that is teachable, auditable, and safe.

Implementation Plan: How to Build the First Owned Interactional DNA Library in 30 Days

This plan is intentionally narrow to keep it doable.

Week 1
Select and instrument workflow

Week 2
Extract data and encode events

Week 3
Operationalise and protect models

Week 4
Run the loop and iterate

Week 1: Select and Instrument

Pick one workflow: support escalation prevention with status update plus exception handling. Define what "good" and "bad" outcomes look like. Set up capture: collect a representative set of recorded calls. Define where the DNA artefacts will live and who owns them.

Privacy requirement: the personalised agents paper highlights privacy and data ownership concerns, emphasising transparency and user control as data collection expands. [8] Treat that as a design constraint, not a legal afterthought.

Week 3: Operationalise and Protect

Embed DNA units into support scripts, manager checklists, and QA routines. Add versioning and change review. Apply secure tool integration practices like least privilege, validation, encryption, and audit trails, aligned with the AWS agent guide's security best practices for tool integration. [13]

Security requirement: include at least one security-focused review, informed by the voice assistant security survey's taxonomy and metrics such as ASR. [10]

1 2 3 4

Week 2: Extract and Encode

Hold three DNA extraction sessions with cross-functional reviewers. Produce 10 DNA units using the schema in Page 9. For each unit, include: trigger, move, cue constraint, constraint capture, boundary rule.

Voice requirement: the persuasion paper's results show voice feature effects and the importance of pragmatics and sociophonetics. [5] Use that as justification for capturing cue constraints explicitly.

Week 4: Run the Loop

Review 20 new calls. Audit adherence to the DNA units. Identify failure modes and refine. Expand to a second workflow only if the loop is stable.

Risk requirement: treat overconfident output as a failure mode, consistent with the LLM limits paper's framing of pattern completion versus generalisable reasoning limits. [7]

Conclusion

Interactional DNA is a practical answer to a common enterprise problem: calls are rich with intent, voice cues, and constraints, but the organisation usually stores them as media and text rather than extracting reusable behaviour. The research base supports that voice interactions contain pragmatic and sociophonetic features that shape outcomes [5], and that humans respond socially to device voice, including alignment effects tied to expressiveness and personification mechanisms. [6] It also shows that personalisation and convenience introduce privacy and data ownership concerns [2] [8], while authenticity and security risks create new failure modes that can produce severe losses and rapid trust collapse. [9] [10] [12]

Treating Interactional DNA as enterprise IP does not require a legal theory to be useful. It requires an operating discipline:

→ **Capture the moves, cues, constraints, and boundaries**

→ **Encode them as a reusable library**

→ **Govern them like a sensitive operational asset**

→ **Instrument them and improve them through stable loops**

→ **Secure them and make them auditable**

> If you do not do this, you can still buy tools and run dashboards. You will not own the intelligence of your calls. **The tool will.**

Bibliography

Full references for all sources cited in this whitepaper.

Ref	Title	Publisher / Journal
[1]	Dethlefs, Nina; Hastie, Helen; Cuayáhuitl, Heriberto; Yu, Yanchao; Rieser, Verena; Lemon, Oliver. *Information Density and Overlap in Spoken Dialogue.*	Manuscript PDF (publisher not stated in the document).
[2]	*Future of Personalized Conversational Agents.*	Research thesis or long-form research document PDF (publisher not stated in the document).
[3]	Paul, Varghese. *Voice AI Agents: The Technical Backbone of Modern Speech Interfaces.*	Sarcouncil Journal of Multidisciplinary (SARC Publisher). Vol 5, Issue 7, 2025.
[4]	Dzhilavyan, Arutiun; Hovakimyan, Areg. *Voice-to-Voice AI.*	American University of Armenia, BS in Data Science (thesis/project). May 9, 2025.
[5]	Jia, Lin; Hu, Xiayu; Yu, Yifan; Liang, Huigang. *AI Voice Matters: AI's Pragmatics and Sociophonetics in Persuasion.*	Research paper PDF (publisher not stated in the document).
[6]	Cohn, Michelle; Zellou, Georgia. *Expressiveness influences human vocal alignment toward voice-AI.*	INTERSPEECH 2019. International Speech Communication Association (ISCA). 2019.
[7]	Zhang, Zheng. *Comprehension Without Competence: Architectural Limits of LLMs in Symbolic Computation and Reasoning.*	Transactions on Machine Learning Research (TMLR). Published 11/2025.
[8]	(Section within) *Future of Personalized Conversational Agents.* Ethical and social implications section on privacy, surveillance, data ownership, and trust.	Publisher not stated in the document.
[9]	Burgess, Jean; Carlon, Dominique; Doyuran, Elif Buse; et al. *Voice AI and authenticity: Current issues and emerging challenges.*	ADM+S Working Paper Series, Working Paper 012. August 2025.
[10]	Li, Jingjin; Chen, Chao; Rahimi Azghadi, Mostafa; Ghodosi, Hossein; Pan, Lei; Zhang, Jun. *Security and privacy problems in voice assistant applications: A survey.*	Computers & Security (Elsevier). 2023.
[11]	Bentivegna, Francesco. *Artificially voiced intelligences: voice and the myth of AI.*	AI & Society (Springer Nature). 2025.
[12]	Li, Jingjin; Chen, Chao; Rahimi Azghadi, Mostafa; Ghodosi, Hossein; Pan, Lei; Zhang, Jun. *Security and privacy problems in voice assistant applications: A survey.*	Computers & Security (Elsevier). 2023.
[13]	Sempf, Aaron; Samuel, Joshua. *Agentic AI frameworks, protocols, and tools on AWS.*	AWS Prescriptive Guidance, Amazon Web Services (AWS). July 2025.
[14]	Sempf, Aaron; Samuel, Joshua. *Agentic AI frameworks, protocols, and tools on AWS.*	AWS Prescriptive Guidance, Amazon Web Services (AWS). July 2025.